Documentation > FeatureStore 0.10.0 documentation

# Databricks Feature Store

# Python API

# Databricks FeatureStoreClients

class databricks.feature\_store.client.FeatureStoreClient(feature\_store\_uri: Optional[str, None] = None, model\_registry\_uri: Op None] = None)

Bases: object

Client for interacting with the Databricks Feature Store.

create\_table(name: str, primary\_keys: Union[str, List[str]], df: Optional[pyspark.sql.dataframe.DataFrame, None] = None, \*, timestamp\_keys: Union[str, List[str], None] = None, partition\_columns: Union[str, List[str], None] = None, schema: Optional[pyspark.sql.types.StructType, None] = None,  $description: Optional[str, None] = None, tags: Optional[Dict[str, str], None] = None, **kwargs) \rightarrow description: Optional[str, None] = None, tags: Optional[Dict[str, str], None] = None, **kwargs) \rightarrow description: Optional[str, None] = None, tags: Optional[Dict[str, str], None] = None, **kwargs) \rightarrow description: Optional[str, None] = None, **kwargs) \leftar{Optional[str, str], None] = None, **kwargs} \rightar{Optional[str, str], None] = None, **kwargs} \rightar{Optional[str], None] = None, **kwargs} \rightar{Optional[$  $databricks. feature\_store. entities. feature\_table. Feature Table$ 

Create and return a feature table with the given name and primary keys.

The returned feature table has the given name and primary keys. Uses the provided schema or the inferred schema of the provided af. If df is provided, this data will be saved in a Delta table. Supported data types for features are: IntegerType, LongType, FloatType, DoubleType, StringType, BooleanType, DateType, TimestampType, ShortType, ArrayType, MapType, and BinaryType, and DecimalType.

- Parameters: name A feature table name of the form <database\_name>.<table\_name>, for example dev.user\_features.
  - primary\_keys The feature table's primary keys. If multiple columns are required, specify a list of column names, for example ['customer\_id', 'region'].
  - df Data to insert into this feature table. The schema of df will be used as the feature table schema.
  - timestamp\_keys -

Columns containing the event time associated with feature value. Timestamp keys and primary keys of the feature table uniquely identify the feature value for an entity at a point in time.

Experimental: This argument may change or be removed in a future release without warning.

partition\_columns -

Columns used to partition the feature table. If a list is provided, column ordering in the list will be used for partitioning.

When choosing partition columns for your feature table, use columns that do not have a high cardinality. An ideal strategy would be such that you expect data in each partition to be at least 1 GB. The most commonly used partition column is a date.

Additional info: Choosing the right partition columns for Delta tables

- schema Feature table schema. Either schema or df must be provided.
- description Description of the feature table.

Tags to associate with the feature table.

# Note

Available in version >= 0.4.1.

• path (Optional[str]) - Path in a supported filesystem. Defaults to the database location.

 $\textbf{register\_table}(\texttt{''}, delta\_table: str, primary\_keys: Union[str, List[str]], timestamp\_keys: Union[str, List[str], None] = None, description: Optional[str, List[str]], timestamp\_keys: Union[str, List[str]], vist[str]], vist[str], vist[st$  $None] = None, tags: Optional[Dict[str, str], None] = None) \rightarrow databricks.feature\_store.entities.feature\_table.FeatureTable = None = No$ 

Register an existing Delta table as a feature table with the given primary keys.

The returned feature table has the same name as the Delta table.

#### Note

Available in version >= 0.3.8.

- Parameters: name A Delta table name of the form <database\_name>.<table\_name>, for example dev.user\_features. The table must exist in the metastore.
  - primary\_keys The Delta table's primary keys. If multiple columns are required, specify a list of column names, for example ['customer\_id', 'region'].
  - timestamp\_keys Columns containing the event time associated with feature value. Together, the timestamp keys and primary keys uniquely identify the feature value at a point in time.
  - description Description of the feature table.
  - tags -

Tags to associate with the feature table.

#### Note

Available in version >= 0.4.1.

A FeatureTable object. Returns:

 $\texttt{get\_table}(\textit{name: str}) \rightarrow \texttt{databricks.feature\_store.entities.feature\_table.FeatureTable}$ 

Get a feature table's metadata.

name - A feature table name of the form <database\_name>.<table\_name>, for example dev.user\_features. Parameters:

 $drop\_table(name: str) \rightarrow None$ 

#### Note

Experimental: This function may change or be removed in a future release without warning.

Delete the specified feature table. This API also drops the underlying Delta table.

# Note

Available in version >= 0.4.1.

Parameters: name - The feature table name of the form <database\_name>.<table\_name>, for example dev.user\_features.

# Note

Deleting a feature table can lead to unexpected failures in upstream producers and downstream consumers (models, endpoints, and scheduled jobs). You must delete any existing published online stores separately.

 $\textbf{read\_table}(\textit{name: str, **kwargs}) \rightarrow \textit{pyspark.sql.dataframe.DataFrame}$ 

Read the contents of a feature table.

name - A feature table name of the form <database\_name>.<table\_name>, for example dev.user\_features. Parameters:

The feature table contents, or an exception will be raised if this feature table does not exist. Returns:

write table(name: str, df: pyspark.sql.dataframe.DataFrame, mode: str = 'merge', checkpoint\_location: Optional[str, None] = None, trigger: Dict[str, Any] = {'processingTime': '5 seconds'}} → Optional[pyspark.sql.streaming.StreamingQuery, None]

Writes to a feature table.

If the input **DataFrame** is streaming, will create a write stream.

- Parameters: name A feature table name of the form <database\_name>.<table\_name>, for example dev.user\_features. Raises an exception if this feature table does not exist.
  - df Spark DataFrame with feature data. Raises an exception if the schema does not match that of the feature table.
  - mode -

Two supported write modes:

- "overwrite" updates the whole table.
- o "merge" will upsert the rows in df into the feature table. If df contains columns not present in the feature table, these columns will be added as new features.
- checkpoint\_location Sets the Structured Streaming checkpointLocation option. By setting a checkpoint\_location, Spark Structured Streaming will store progress information and intermediate state, enabling recovery after failures. This parameter is only supported when the argument **df** is a streaming **DataFrame**.
- trigger If df.isStreaming, trigger defines the timing of stream data processing, the dictionary will be unpacked and passed to DataStreamWriter.trigger as arguments. For example, trigger={'once': True} will result in a call to DataStreamWriter.trigger(once=True).

If df.isStreaming, returns a PySpark StreamingQuery. None otherwise. Returns:

add data  $sources(*, feature\_table\_name: str, source\_names: Union[str, List[str]], source\_type: str = 'custom') <math>\rightarrow$  None

#### Note

Experimental: This function may change or be removed in a future release without warning.

Add data sources to the feature table.

Parameters:

- feature\_table\_name The feature table name.
- source\_names Data source names. For multiple sources, specify a list. If a data source name already exists, it is ignored.
- · source\_type -

One of the following:

- "table": Table in format <database\_name>.<table\_name> and is stored in the metastore (eg Hive).
- o "path": Path, eg in the Databricks File System (DBFS).
- "custom": Manually added data source, neither a table nor a path.

 $\textbf{delete\_data\_sources}(\,{}^\star, \textit{feature\_table\_name: str, source\_names: Union[str, List[str]]}) \rightarrow \textbf{None}$ 

#### Note

Experimental: This function may change or be removed in a future release without warning.

Delete data sources from the feature table.

#### Note

Data sources of all types (table, path, custom) that match the source names will be deleted.

Parameters:

- feature table name The feature table name.
- source\_names Data source names. For multiple sources, specify a list. If a data source name does not exist, it is ignored.

publish\_table(name: str, online\_store: databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpec, \*, filter\_condition: Optional[str, None] = None, mode: str = 'merge', streaming: bool = False, checkpoint\_location: Optional[str, None] = None, trigger: Dict[str, Any] = Publish a feature table to an online store.

- Parameters: name Name of the feature table.
  - online\_store Specification of the online store.
  - filter\_condition A SQL expression using feature table columns that filters feature rows prior to publishing to the online store. For example, "dt > '2020-09-10'". This is analogous to running df.filter or a WHERE condition in SQL on a feature table prior to publishing.
  - mode -

Specifies the behavior when data already exists in this feature table in the online store. If "overwrite" mode is used, existing data is replaced by the new data. If "merge" mode is used, the new data will be merged in, under these conditions:

- o If a key exists in the online table but not the offline table, the row in the online table is unmodified.
- o If a key exists in the offline table but not the online table, the offline table row is inserted into the online table.
- o If a key exists in both the offline and the online tables, the online table row will be updated.
- streaming If True, streams data to the online store.
- checkpoint\_location Sets the Structured Streaming checkpointLocation option. By setting a checkpoint\_location, Spark Structured Streaming will store progress information and intermediate state, enabling recovery after failures. This parameter is only supported when **streaming=True**.
- trigger If streaming=True, trigger defines the timing of stream data processing. The dictionary will be unpacked and passed to DataStreamWriter.trigger as arguments. For example, trigger={'once': True} will result in a call to DataStreamWriter.trigger(once=True).
- features –

Specifies the feature column(s) to be published to the online store. The selected features must be a superset of existing online store features. Primary key columns and timestamp key columns will always be published.

#### Note

This parameter is only supported when mode="merge". When features is not set, the whole feature table will be published.

If streaming=True, returns a PySpark StreamingQuery, None otherwise. Returns:

create training set(df: pyspark.sql.dataframe.DataFrame, feature\_lookups: List[databricks.feature\_store.entities.feature\_lookup.FeatureLookup],  $label: \textit{Union[str, List[str], None]}, \textit{exclude\_columns: Optional[List[str], None]} = \textit{None}) \rightarrow \textit{databricks.feature\_store.training\_set.TrainingSet}$ 

Create a **TrainingSet**.

- Parameters: df The DataFrame used to join features into.
  - feature\_lookups List of features to join into the DataFrame.
  - label Names of column(s) in DataFrame that contain training set labels. To create a training set without a label field, i.e. for unsupervised training set, specify label = None.
  - exclude\_columns Names of the columns to drop from the TrainingSet DataFrame.

A TrainingSet object. Returns:

log model(model: Any, artifact\_path: str, \*, flavor: module, training\_set: Optional[databricks.feature\_store.training\_set.TrainingSet, None] = None, registered\_model\_name: Optional[str, None] = None, await\_registration\_for: int = 300, \*\*kwargs)

Log an MLflow model packaged with feature lookup information.

# Note

The DataFrame returned by TrainingSet.load df() must be used to train the model. If it has been modified (for example data normalization, add a column, and similar), these modifications will not be applied at inference time, leading to training-serving skew.

- Parameters: model Model to be saved. This model must be capable of being saved by flavor.save\_model. See the MLflow Model API.
  - artifact path Run-relative artifact path.
  - flavor MLflow module to use to log the model. flavor should have type ModuleType. The module must have a method save\_model, and must support the python\_function flavor. For example, mlflow.sklearn, mlflow.xgboost, and similar.
  - training\_set The TrainingSet used to train this model.
  - registered\_model\_name -

#### Note

Experimental: This argument may change or be removed in a future release without warning.

If given, create a model version under registered\_model\_name, also creating a registered model if one with the given name does

• await\_registration\_for - Number of seconds to wait for the model version to finish being created and is in READY status. By default, the function waits for five minutes. Specify o or None to skip waiting.

None Returns:

 $score\_batch(model\_uri: str, df: pyspark.sql.dataframe.DataFrame, result\_type: str = 'double') <math>\rightarrow$  pyspark.sql.dataframe.DataFrame

Evaluate the model on the provided **DataFrame**.

Additional features required for model evaluation will be automatically retrieved from Feature Store.

The model must have been logged with FeatureStoreClient.log model(), which packages the model with feature metadata. Unless present in df, these features will be looked up from Feature Store and joined with df prior to scoring the model.

If a feature is included in df, the provided feature values will be used rather than those stored in Feature Store.

For example, if a model is trained on two features account creation date and num lifetime purchases, as in:

```
feature lookups = [
   FeatureLookup(
       table_name = 'trust_and_safety.customer_features',
       feature_name = 'account_creation_date',
       lookup_key = 'customer_id',
    ).
    FeatureLookup(
        table name = 'trust and safety.customer features',
        feature name = 'num lifetime purchases',
        lookup_key = 'customer_id'
1
with mlflow.start_run():
    training_set = fs.create_training_set(
       df.
       feature_lookups = feature_lookups,
       label = 'is banned',
       exclude_columns = ['customer_id']
     fs.log_model(
       model.
       "model",
       flavor=mlflow.sklearn,
       training_set=training_set,
       registered_model_name="example_model"
```

Then at inference time, the caller of FeatureStoreClient.score\_batch() must pass a DataFrame that includes customer\_id, the lookup\_key specified in the FeatureLookups of the training\_set. If the DataFrame contains a column account\_creation\_date, the values of this column will be used in lieu of those in Feature Store. As in:

```
# batch_df has columns ['customer_id', 'account_creation_date']
predictions = fs.score_batch(
    'models:/example_model/1',
   batch_df
```

Parameters: • model\_uri -

The location, in URI format, of the MLflow model logged using FeatureStoreClient.log\_model(). One of:

- o runs:/<mlflow\_run\_id>/run-relative/path/to/model
- o models:/<model\_name>/<model\_version>
- o models:/<model\_name>/<stage>

For more information about URI schemes, see Referencing Artifacts.

The DataFrame to score the model on. Feature Store features will be joined with df prior to scoring the model. df must:

- 1. Contain columns for lookup keys required to join feature data from Feature Store, as specified in the feature\_spec.yaml artifact.
- 2. Contain columns for all source keys required to score the model, as specified in the feature\_spec.yaml artifact.
- 3. Not contain a column prediction, which is reserved for the model's predictions. df may contain additional columns.
- result\_type The return type of the model. See mlflow.pyfunc.spark\_udf() result\_type.

Returns:

A DataFrame containing:

- 1. All columns of df.
- 2. All feature values retrieved from Feature Store.
- 3. A column **prediction** containing the output of the model.

 $\verb|set_feature_table_tag|(*, table_name: str, key: str, value: str)| \rightarrow \verb|None|$ 

Create or update a tag associated with the feature table. If the tag with the corresponding key already exists, its value will be overwritten with the new value.

Note

Available in version >= 0.4.1.

Parameters:

- table\_name the feature table name
- key tag key
- value tag value

 $delete\_feature\_table\_tag(*, table\_name: str, key: str) \rightarrow None$ 

Delete the tag associated with the feature table. Deleting a non-existent tag will emit a warning.

Note

Available in version >= 0.4.1.

Parameters:

- table\_name the feature table name.
- key the tag key to delete.

create\_feature\_serving\_endpoint(\*, endpoint\_name: str, feature\_lookups:  $\textit{List[databricks.feature\_store.entities.feature\_lookup.FeatureLookup])} \rightarrow$  $databricks. feature\_store. entities. feature\_serving\_endpoint. Feature Serving Endpoint$ 



Experimental: This function may change or be removed in a future release without warning.

Experimental feature

 $\label{eq:get_feature_serving_endpoint} \textbf{(*, endpoint\_name)} \rightarrow \textbf{databricks.feature\_store.entities.feature\_serving\_endpoint.FeatureServingEndpoint}$ 

Note

Experimental: This function may change or be removed in a future release without warning.

Experimental feature

delete\_feature\_serving\_endpoint(\*, endpoint\_name) → None

Note

Experimental: This function may change or be removed in a future release without warning.

Experimental feature

# Feature Lookup

class databricks.feature\_store.entities.feature\_lookup.FeatureLookup(table\_name: str, lookup\_key: Union[str, List[str]], \*, feature\_names: Union[str, List[str], None] = None, rename\_outputs: Optional[Dict[str, str], None] = None, timestamp\_lookup\_key: Union[str, List[str], None] =

Bases: databricks.feature\_store.entities.\_feature\_store\_object.\_FeatureStoreObject

Value class used to specify a feature to use in a **TrainingSet**.

- Parameters: table\_name Feature table name.
  - lookup\_key Key to use when joining this feature table with the DataFrame passed to FeatureStoreClient.create\_training\_set(). The lookup\_key must be the columns in the DataFrame passed to FeatureStoreClient.create training set(). The type and order of lookup key columns in that DataFrame must match the primary key of the feature table referenced in this FeatureLookup.
  - feature\_names A single feature name, a list of feature names, or None to lookup all features (excluding primary keys) in the feature table at the time that the training set is created. If your model requires primary keys as features, you can declare them as independent FeatureLookups.
  - rename\_outputs If provided, renames features in the TrainingSet returned by of FeatureStoreClient.create\_training\_set.
  - timestamp lookup kev -

Key to use when performing point-in-time lookup on this feature table with the DataFrame passed to  $\textbf{FeatureStoreClient.create\_training\_set}() . The \ \textbf{timestamp\_lookup\_key} \ \text{must} \ \text{be} \ \text{the} \ \text{columns} \ \text{in} \ \text{the} \ \text{DataFrame} \ \text{passed} \ \text{to} \ \text{the} \ \text{$ FeatureStoreClient.create\_training\_set(). The type of timestamp\_lookup\_key columns in that DataFrame must match the type of the timestamp key of the feature table referenced in this FeatureLookup.

Note

Experimental: This argument may change or be removed in a future release without warning.

- feature\_name Feature name. Deprecated as of version 0.3.4. Use feature names.
- output\_name If provided, rename this feature in the output of FeatureStoreClient.create\_training\_set.Deprecated as of version 0.3.4. Use rename\_outputs.

init (table\_name: str, lookup\_key: Union[str, List[str]], \*, feature\_names: Union[str, List[str], None] = None, rename\_outputs: Optional[Dict[str, str], None] = None, timestamp\_lookup\_key: Union[str, List[str], None] = None, \*\*kwargs)

Initialize a FeatureLookup object.

#### table name

The table name to use in this FeatureLookup.

The lookup key(s) to use in this FeatureLookup.

#### feature\_name

The feature name to use in this FeatureLookup. Deprecated as of version 0.3.4. Use feature names.

#### output name

The output name to use in this FeatureLookup. **Deprecated** as of version 0.3.4. Use **feature\_names**.

# **Training Set**

class databricks.feature\_store.training\_set.TrainingSet(feature\_spec: databricks.feature\_store.entities.feature\_spec.FeatureSpec, df: pyspark.sql.dataframe.DataFrame, labels: List[str], feature\_table\_metadata\_map: Dict[str, databricks.feature\_store.entities.feature\_table.FeatureTable], feature\_table\_data\_map: Dict[str, pyspark.sql.dataframe.DataFrame])

Bases: object

Class that defines **TrainingSet** objects.

#### Note

The TrainingSet constructor should not be called directly. Instead, call FeatureStoreClient.create training set.

#### **load\_df**() → pyspark.sql.dataframe.DataFrame

Load a DataFrame.

Return a DataFrame for training.

The returned **DataFrame** has columns specified in the **feature\_spec** and **labels** parameters provided in **FeatureStoreClient.create\_training\_set**.

Returns:

A DataFrame for training

# Feature Table

#### Classes

class databricks.feature\_store.entities.feature\_table.FeatureTable(name, table\_id, description, primary\_keys, partition\_columns, features, creation\_timestamp=None, online\_stores=None, notebook\_producers=None, job\_producers=None, table\_data\_sources=None, path\_data\_sources=None, custom\_data\_sources=None, timestamp\_keys=None, tags=None)

Value class describing one feature table.

This will typically not be instantiated directly, instead the FeatureStoreClient.create table will create FeatureTable objects.

# Online Store Spec

class databricks.feature\_store.online\_store\_spec.AmazonRdsMySqlSpec(hostname: str, port: int, user: Optional[str, None] = None, password: Optional[str, None] = None, database\_name: Optional[str, None] = None, table\_name: Optional[str, None] = None, driver\_name: Optional[str, None] = None, read\_secret\_prefix: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None)

Bases: databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpec

Class that defines and creates AmazonRdsMySqlSpec objects.

This OnlineStoreSpec implementation is intended for publishing features to Amazon RDS MySQL and Aurora (MySQL-compatible edition).

See OnlineStoreSpec documentation for more usage information, including parameter descriptions.

#### Parameters:

- hostname Hostname to access online store.
- port Port number to access online store.
- user Username that has access to the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
- password Password to access the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
- database\_name Database name.
- table\_name Table name.
- driver\_name Name of custom JDBC driver to access the online store.
- read\_secret\_prefix Prefix for read secret.
- write\_secret\_prefix Prefix for write secret.

#### hostname

Hostname to access the online store.

#### port

Port number to access the online store.

#### database name

Database name.

#### cloud

Define the cloud propert for the data store.

#### store\_type

Define the data store type property.

#### auth\_type()

Publish Auth type.

class databricks.feature\_store.online\_store\_spec.AzureMySqlSpec(hostname: str, port: int, user: Optional[str, None] = None, password: Optional[str, None] = None, database\_name: Optional[str, None] = None, table\_name: Optional[str, None] = None, driver\_name: Optional[str, None] = None, read\_secret\_prefix: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None)

Bases: databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpec

Define the AzureMySqlSpec class.

This OnlineStoreSpec implementation is intended for publishing features to Azure Database for MySQL.

See OnlineStoreSpec documentation for more usage information, including parameter descriptions.

#### Parameters:

- hostname Hostname to access online store.
- port Port number to access online store.
- user Username that has access to the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
- password Password to access the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
- database\_name Database name.
- table\_name Table name.
- driver\_name Name of custom JDBC driver to access the online store.
- read\_secret\_prefix Prefix for read secret.
- write\_secret\_prefix Prefix for write secret.

#### hostname

Hostname to access the online store.

#### port

Port number to access the online store.

#### database\_name

Database name.

#### cloud

Define the cloud the fature store runs.

# store\_type

Define the data store type.

#### auth\_type()

Publish Auth type.

class databricks.feature\_store.online\_store\_spec.AzureSqlServerSpec(hostname: str, port: int, user: Optional[str, None] = None, password: Optional[str, None] = None, database\_name: Optional[str, None] = None, table\_name: Optional[str, None] = None, driver\_name: Optional[str, None] = None, read\_secret\_prefix: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None)

Bases: databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpec

This **onlineStoreSpec** implementation is intended for publishing features to Azure SQL Database (SQL Server).

The spec supports SQL Server 2019 and newer.

See OnlineStoreSpec documentation for more usage information, including parameter descriptions.

#### Parameters:

- hostname Hostname to access online store.
- port Port number to access online store.
- user Username that has access to the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
- password Password to access the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
- database\_name Database name.
- table name Table name.
- driver\_name Name of custom JDBC driver to access the online store.
- read\_secret\_prefix Prefix for read secret.
- write\_secret\_prefix Prefix for write secret.

#### hostname

Hostname to access the online store.

## port

Port number to access the online store.

# database\_name

Database name.

#### cloud

Define the cloud the fature store runs.

# store\_type

Define the data store type.

### auth\_type()

Publish Auth type.

class databricks.feature\_store.online\_store\_spec.AmazonDynamoDBSpec(\*, region: Optional[str, None], access\_key\_id: Optional[str, None] = None, secret\_access\_key: Optional[str, None] = None, session\_token: Optional[str, None] = None, table\_name: Optional[str, None] = None, read\_secret\_prefix: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None, table\_name: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None, table\_name: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None, table\_name: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None, table\_name: Optional[str, None] = None, table\_nam

Bases: databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpec

This OnlineStoreSpec implementation is intended for publishing features to Amazon DynamoDB.

If table\_name is not provided, FeatureStoreClient.publish\_table will use the offline store's database and table name combined as the online table name.

To use a different table name in the online store, provide a value for the table\_name argument.

The expected read or write secrets for DynamoDB for a given {prefix} string are \${prefix}-access-key-id,

\${prefix}-secret-access-key, and \${prefix}-session-token.

If none of the access\_key\_id, secret\_access\_key, and write\_secret\_prefix are passed, the instance profile attached to the cluster will be used to write to DynamoDB.

#### Note

AmazonDynamoDBSpec is available in version >= 0.3.8.

Instance profile based writes are available in version >= 0.4.1.

- **Parameters:** region Region to access online store.
  - access\_key\_id Access key ID that has access to the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
  - secret\_access\_key Secret access key to access the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
  - session\_token Session token to access the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
  - table\_name Table name.
  - read\_secret\_prefix Prefix for read secret.
  - write\_secret\_prefix Prefix for write secret.
  - ttl The time to live for data published to the online store. This attribute is only applicable when publishing time series feature tables. If the time to live is specified for a time series table, FeatureStoreClient.publish\_table() will publish a window of data instead of the latest snapshot.

# access\_key\_id

# **A** Warning

databricks.feature store.online\_store\_spec.amazon\_dynamodb\_online\_store\_spec.AmazonDynamoDBSpec.access\_key\_id is deprecated since v0.6.0. This method will be removed in a future release. Use write secret prefix instead.

Access key ID that has access to the online store. Property will be empty if write\_secret\_prefix or the instance profile attached to the cluster are intended to be used.

### secret\_access\_key

### **A** Warning

databricks.feature store.online store spec.amazon dynamodb online store spec.AmazonDynamoDBSpec.secret access key is deprecated since v0.6.0. This method will be removed in a future release. Use write\_secret\_prefix instead.

Secret access key to access the online store. Property will be empty if write\_secret\_prefix or the instance profile attached to the cluster are intended to be used.

## session\_token

### **A** Warning

databricks.feature\_store.online\_store\_spec.amazon\_dynamodb\_online\_store\_spec.AmazonDynamoDBSpec.session\_token is deprecated since v0.6.0. This method will be removed in a future release. Use write\_secret\_prefix instead.

Session token to access the online store. Property will be empty if write\_secret\_prefix or the instance profile attached to the cluster are intended to be used.

#### cloud

Define the cloud property for the data store.

#### store\_type

Define the data store type.

# region

Region to access the online store.

#### ttl

Time to live attribute for the online store.

#### auth\_type()

Publish Auth type.

class databricks.feature\_store.online\_store\_spec.AzureCosmosDBSpec(\*, account\_uri: str, database\_name: Optional[str, None] = None, container\_name: Optional[str, None] = None, read\_secret\_prefix: Optional[str, None] = None, write\_secret\_prefix: str)

 ${\tt Bases: databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpecsure\_specsure\_$ 

This OnlineStoreSpec implementation is intended for publishing features to Azure Cosmos DB.

If database\_name and container\_name are not provided, FeatureStoreClient.publish\_table will use the offline store's database and table name as the Cosmos DB database and container name.

The expected read or write secret for Cosmos DB for a given {prefix} string is \${prefix}-authorization-key.

The authorization key can be either the Cosmos DB account primary or secondary key.

#### Note

Available in version >= 0.5.0.

#### Parameters:

- account\_uri URI of the Cosmos DB account.
- database\_name Database name.
- container\_name Container name.
- read\_secret\_prefix Prefix for read secret.
- write\_secret\_prefix Prefix for write secret.

# account\_uri

Account URI of the online store.

# database\_name

Database name.

#### container name

Container name.

#### cloud

Define the cloud property for the data store.

#### store\_type

Define the data store type.

### auth\_type()

Publish Auth type.

class databricks.feature\_store.online\_store\_spec.OnlineStoreSpec(\_type, hostname: [<class 'str'>, None] = None, port: [<class 'int'>, None] = None, user: Optional[str, None] = None, password: Optional[str, None] = None, database\_name: Optional[str, None] = None, table\_name: Optional[str, None] = None, driver\_name: Optional[str, None] = None, read\_secret\_prefix: Optional[str, None] = None, write\_secret\_prefix: Optional[str, None] = None, \_internal\_properties: Optional[Dict[str, str], None] = None)

Bases: abc.ABC

Parent class for all types of **OnlineStoreSpec** objects.

Abstract base class for classes that specify the online store to publish to.

If database\_name and table\_name are not provided, FeatureStoreClient.publish\_table will use the offline store's database and table names.

To use a different database and table name in the online store, provide values for both database name and table name arguments.

The JDBC driver can be customized with the optional driver name argument. Otherwise, a default is used.

Strings in the primary key should not exceed 100 characters.

The online database should already exist.

#### Note

It is strongly suggested (but not required), to provide read-only database credentials via the read secret prefix in order to grant the least amount of database access privileges to the served model. When providing a read\_secret\_prefix, the secrets must exist in the scope name using the expected format, otherwise **publish** table will return an error.

- Parameters: hostname Hostname to access online store. The database hostname cannot be changed. Subsequent publish calls to the same online store must provide the same hostname.
  - port Port number to access online store. The database port cannot be changed. Subsequent publish calls to the same online store must provide the same port.
  - user Username that has write access to the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
  - password Password to access the online store. Deprecated as of version 0.6.0. Use write\_secret\_prefix instead.
  - database\_name Database name.
  - table\_name Table name.
  - driver\_name Name of custom JDBC driver to access the online store.
  - read\_secret\_prefix -

The secret scope name and secret key name prefix where read-only online store credentials are stored. These credentials will be used during online feature serving to connect to the online store from the served model. The format of this parameter should be \${scope-name}/\${prefix}, which is the name of the secret scope, followed by a /, followed by the secret key name prefix. The scope passed in must contain the following keys and corresponding values:

- \${prefix}-user where \${prefix} is the value passed into this function. For example if this function is called with datascience/staging, the datascience secret scope should contain the secret named staging-user, which points to a secret value with the database username for the online store.
- \${prefix}-password where \${prefix} is the value passed into this function. For example if this function is called with datascience/staging, the datascience secret scope should contain the secret named staging-password, which points to a secret value with the database password for the online store.

Once the read\_secret\_prefix is set for an online store, it cannot be changed.

#### write secret prefix -

The secret scope name and secret key name prefix where read-write online store credentials are stored. These credentials will be used to connect to the online store to publish features. If user and password are passed, this field must be None, or an exception will be raised. The format of this parameter should be \${scope-name}/\${prefix}, which is the name of the secret scope, followed by a /, followed by the secret key name prefix. The scope passed in must contain the following keys and corresponding values:

- \${prefix}-user where \${prefix} is the value passed into this function. For example if this function is called with datascience/staging, the datascience secret scope should contain the secret named staging-user, which points to a secret value with the database username for the online store.
- \${prefix}-password where \${prefix} is the value passed into this function. For example if this function is called with datascience/staging, the datascience secret scope should contain the secret named staging-password, which points to a secret value with the database password for the online store.

# type

Type of the online store.

# table\_name

Table name.

#### user

databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpec.user is deprecated since v0.6.0. This method will be removed in a future release. Use write\_secret\_prefix instead.

Username that has access to the online store.

Property will be empty if write\_secret\_prefix argument was used.

# password

#### **A** Warning

databricks.feature\_store.online\_store\_spec.online\_store\_spec.OnlineStoreSpec.password is deprecated since v0.6.0. This method will be removed in a future release. Use write\_secret\_prefix instead.

Password to access the online store.

Property will be empty if write secret prefix argument was used.

#### driver

Name of the custom JDBC driver to access the online store.

#### read\_secret\_prefix

Prefix for read access to online store.

Name of the secret scope and prefix that contains the username and password to access the online store with read-only credentials.

See the **read\_secret\_prefix** parameter description for details.

# write\_secret\_prefix

Secret prefix that contains online store login info.

Name of the secret scope and prefix that contains the username and password to access the online store with read/write credentials. See the write\_secret\_prefix parameter description for details.

#### cloud

Cloud provider where this online store is located.

# store\_type

Store type.

# auth\_type()

Publish Auth type.